SIMD implementation

Example: exclusive scan 32-element array

32-wide GPU execution (SPMD program)

Exclusive result is returned (note: state of ptr[] is inclusive result)

CUDA thread index

```
template<class OP, class T>
__device__ T scan_warp(volatile T *ptr, const unsigned int idx)
{
    const unsigned int lane = idx & 31; // index of thread in warp (0..31)

    if (lane >= 1)  ptr[idx] = OP::apply(ptr[idx - 1],  ptr[idx]);
    if (lane >= 2)  ptr[idx] = OP::apply(ptr[idx - 2],  ptr[idx]);
    if (lane >= 4)  ptr[idx] = OP::apply(ptr[idx - 4],  ptr[idx]);
    if (lane >= 8)  ptr[idx] = OP::apply(ptr[idx - 8],  ptr[idx]);
    if (lane >= 16) ptr[idx] = OP::apply(ptr[idx - 16], ptr[idx]);

    return (lane>0) ? ptr[idx-1] : OP::identity();
}
```

Work: ??